

# Optimization of RDMA communication on an Infiniband network

Presentation no .1

Stamoulis Spyros 03775

Professors:

- Christos Antonopoulos
- Athanasios Fevgas

# Initial steps

---

- Studied the samples
  - rdma\_send/rdma\_receive
  - rdma\_write,rdma\_read
- After I understood the functionality of the code I tried to generalize the send receive sample to send many messages.

# Version 1

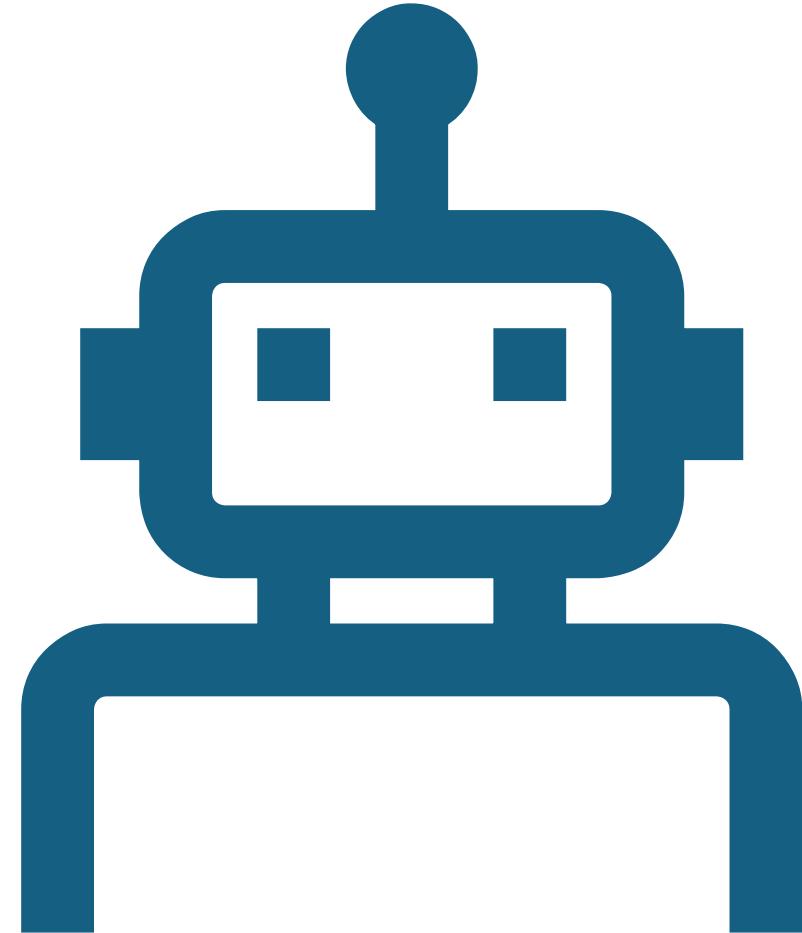
---

- As a first change I changed the rdma\_send/rdma\_receive from a single message program to send/receive 1000 messages and benchmarking the time.

# Goal

---

- **Original (Before):** The program was designed to:
  - Connect to the receiver.
  - Prepare *one* buffer.
  - Submit *one* send task.
  - Clean up and exit.
- **New Goal (After):** The benchmark needs to:
  - Connect to the receiver.
  - Start a timer.
  - Send packets.
  - Stop the timer.
  - Calculate and report the throughput (Gbps).



Send sample

# Filling the *Entire* Buffer

---

- To measure throughput, we must send the full buffer size, not just a small string.
- I replaced `strncpy` (which copies a short string) with `memset`.
- This ensures we are benchmarking the full `MAX_BUFF_SIZE` (e.g., 1MB) on every send.

```
/* Set data of src buffer */
result = doca_buf_get_data(resources->src_buf, &src_buf_data);
if (result != DOCA_SUCCESS) {
    DOCA_LOG_ERR("Failed to get source buffer data: %s", doca_error_get_descr(result));
    goto destroy_src_buf;
}
DOCA_LOG_INFO("did doca_buf_get_data\n");
// strncpy(src_buf_data, resources->cfg->send_string, MAX_BUFF_SIZE + 1);
//fill the string with data
memset(src_buf_data, 'S', MAX_BUFF_SIZE);

DOCA_LOG_INFO("did memset\n");
```

The image on the right is from the func prepare and submit task:

# Starting the Benchmark

- Added logic to rdma\_send\_prepare\_and\_submit\_task to start the benchmark.
- It uses a new counter, transfers\_left, to track the state.
- On the very first call (transfers\_left == 0), it starts the high-resolution timer.

```
if(resources->transfers_left == 0){  
    DOCA_LOG_INFO("Starting benchmark: %d transfers of %d bytes", NUM_TRANSFERS, MAX_BUFF_SIZE);  
    resources->transfers_left = NUM_TRANSFERS;  
    clock_gettime(CLOCK_MONOTONIC, &resources->start_time);  
  
    if (resources->cfg->use_rdma_cm == true) {  
        DOCA_LOG_INFO("Please press enter after the receive task has been successfully submitted in the rec  
  
        /* Wait for enter */  
        wait_for_enter();  
    }  
}
```

# The Benchmark Logic

- The completion callback (`rdma_send_completed_callback`) became the heart of the new version.
- When a send completes, the callback does the following things:
  1. Decrements the `transfers_left` counter.
  2. If `transfers_left > 0`, it immediately calls `rdma_send_prepare_and_submit_task` again.
  3. Ends the benchmark if there are no more tasks to do.

```
else{
    //print benchmark results
    clock_gettime(CLOCK_MONOTONIC, &resources->end_time);
    double total_bytes = (double)NUM_TRANSFERS * MAX_BUFF_SIZE;
    double elapsed_seconds = (resources->end_time.tv_sec - resources->start_time.tv_sec) +
    (resources->end_time.tv_nsec - resources->start_time.tv_nsec) / 1e9;
    double gbps = (total_bytes * 8) / (elapsed_seconds * 1e9);

    printf("\n--- BENCHMARK RESULTS ---\n");
    printf("Total data sent: %.2f MB\n", total_bytes / (1024*1024));
    printf("Elapsed time:    %.4f seconds\n", elapsed_seconds);
    printf("Throughput:     %.4f Gbps\n", gbps);
    printf("-----\n\n");

    if (resources->cfg->use_rdma_cm == true){
        (void)rdma_cm_disconnect(resources);
    }
    (void)doca_ctx_stop(resources->rdma_ctx);
}
```

```
//reduce number of tasks to run
resources->transfers_left--;

resources->num_remaining_tasks--;

if (resources->transfers_left > 0){
    result = rdma_send_prepare_and_submit_task(resources);
    if (result != DOCA_SUCCESS){
        DOCA_LOG_ERR("rdma_send_prepare_and_submit_task() failed: %s", doca_error_get_descr(result));
        if (resources->cfg->use_rdma_cm == true){
            (void)rdma_cm_disconnect(resources);
        }
        (void)doca_ctx_stop(resources->rdma_ctx);
    }
} else{
```

# Receive sample

---

- On the receive side I did almost the same changes to be able to receive `NUM_TRANSFERS` number of messages

# Version 2

---

- Here I tried to implement a pipeline of the tasks to improve the throughput and I experimented with different pipe sizes
- I only changed the send file

# In the rdma\_send\_state\_change\_callback



I submit PIPELINE DEPTH tasks instead of just one in the beginning

```
case DOCA_CTX_STATE_RUNNING:
    DOCA_LOG_INFO("RDMA context is running");

    result = rdma_send_export_and_connect(resources);
    if (result != DOCA_SUCCESS) {
        DOCA_LOG_ERR("rdma_send_export_and_connect() failed: %s", doca_error_get_descr(result));
        break;
    } else
        DOCA_LOG_INFO("RDMA context finished initialization");

    if (cfg->use_rdma_cm == true)
        break;
    // result = rdma_send_prepare_and_submit_task(resources);
    for (int i = 0; i < PIPELINE_DEPTH; i++) {
        result = rdma_send_prepare_and_submit_task(resources);
        if (result != DOCA_SUCCESS){
            DOCA_LOG_ERR("rdma_send_prepare_and_submit_task() failed: %s", doca_error_get_descr(result));
            }
    }
    break;
```

# Version 3

---

- I moved to the write sample and I did the same changes in the requester sample as the ones I did at version 1 on the send sample in order to make possible the multiple write tasks.
- This sample was waiting for the user to press enter to finish so I added a send task to let the receiver (write\_responder) that the writing has finished. Before the send task, the sender (write\_requester) stops the benchmark and prints the results.
- Finally when the receiver gets the final send task it stops the context.

# Requester (client)

- Added send task when the benchmark is finished
- Also added standard callbacks to handle the ending of the send task.

```
DOCA_LOG_INFO("All write tasks completed. Sending final signal to responder.");
struct doca_rdma_task_send *send_task;
struct doca_buf *signal_buf;
void *buf_data;
const char *signal_msg = "done";
// 1. Get a buffer from the inventory for the signal
result = doca_buf_inventory_buf_get_by_data(resources->buf_inventory, resources->mmap,
                                              resources->mmap_memrange+SIGNAL_OFFSET, strlen(signal_msg),
                                              &signal_buf);
if (result != DOCA_SUCCESS) {
    DOCA_LOG_ERR("Failed to get buffer for signal task: %s", doca_error_get_descr(result));
    (void)doca_ctx_stop(resources->rdma_ctx);
    return;
}
doa_buf_get_data(signal_buf, &buf_data);
strcpy(buf_data, signal_msg);
result = doca_rdma_task_send_allocate_init(resources->rdma, resources->connections[0],
                                             signal_buf, // Pass the doca_buf* here
                                             (union doca_data){0}, &send_task);
if (result == DOCA_SUCCESS) {
    doca_task_submit(doca_rdma_task_send_as_task(send_task));
} else {
    DOCA_LOG_ERR("Failed to allocate signal task: %s", doca_error_get_descr(result));
    doca_buf_dec_refcount(signal_buf, NULL); // Clean up the buffer if task allocation fails
}
```

# Responder (server)

- Added final\_signal\_received\_callback handles when the receive task is completed and checks 64 bytes in the memory to make sure the write operation was correct.
- I also added a standard error callback func that I copied from the receive sample
- And in similar fashion a post receive task func

```
static void final_signal_received_callback(struct doca_rdma_task_receive *task,
                                         union doca_data task_user_data,
                                         union doca_data ctx_user_data)
{
    doca_error_t result = DOCA_SUCCESS;
    // char buffer[MAX_BUFF_SIZE+1];
    struct rdma_resources *resources = (struct rdma_resources *)ctx_user_data.ptr;
    DOCA_LOG_INFO("\n Received final signal from requester. Benchmark complete.\n");

    /* The RDMA Write target memory (where requester wrote data) */
    char *rdma_written_data = (char *)resources->mmap_memrange;

    /* Print the first bytes written by the requester */
    char print_buf[65];
    memcpy(print_buf, rdma_written_data, 64);
    print_buf[64] = '\0';

    DOCA_LOG_INFO("First 64 bytes written via RDMA Write: \"%s\"", print_buf);

    /* Print the received message ("done") separately */
    struct doca_buf *recv_buf = doca_rdma_task_receive_get_dst_buf(task);
    void *recv_data = NULL;
    size_t recv_len = 0;
    doca_buf_get_data(recv_buf, &recv_data);

    DOCA_LOG_INFO("Final signal message: \"%s\"", (char *)recv_data);
    // Free the task and its buffer
    doca_task_free(doca_rdma_task_receive_as_task(task));
    result = doca_buf_dec_refcount(resources->dst_buf, NULL);
    if (result != DOCA_SUCCESS) {
        DOCA_LOG_ERR("Failed to decrease src buf count: %d", doca_error_get_desc(result));
    }
}
```

Also in the state change func I changed the function called. So instead of actively waiting it waits for a receive task to finish

```
case DOCA_CTX_STATE_RUNNING:
    DOCA_LOG_INFO("RDMA context is running");

    result = rdma_write_responder_export_and_connect(resources);
    if (result != DOCA_SUCCESS) {
        DOCA_LOG_ERR("rdma_write_responder_export_and_connect() failed: %s",
                     doca_error_get_descr(result));
        break;
    } else
        DOCA_LOG_INFO("RDMA context finished initialization");

    if (cfg->use_rdma_cm == true)
        break;

// result = responder_wait_for_requester_finish(resources);
result = post_final_signal_receive(resources);
if (result != DOCA_SUCCESS) {
    DOCA_LOG_ERR("Failed to post final signal receive task: %s", doca_error_get_descr(result));
    (void)doctx_stop(ctx);
}
break;
```

# Times

---

- Note:
  - For all tests total data sent= 62.50 MB.
  - For the previous versions, the throughput is negligible (near zero)
- Times:
  - Throughput: 5.1881 Gbps
  - Throughput: 2.8214 Gbps
  - Throughput: 5.1820 Gbps
  - Throughput: 3.8275 Gbps
  - Throughput: 5.1470 Gbps
- Average = 4.433200
- Standard deviation = 1.072947

# Version 4 (Also fixed a few issues)

---

- I added pipeline and I exported only once the mmap instead of doing it in every task submission.
- Each message overwrites the previous one
- Issue with the freeing of the buffers
- When the 1000<sup>th</sup> task is submitted the benchmark ends. So we don't wait for it to finish

# Requester

---

I added a function to fill the pipeline in the beginning and export the mmap

```
static doca_error_t cm_requester_setup_and_submit(struct rdma_resources *resources)
{
    doca_error_t result;

    result = doca_mmap_create_from_export(NULL,
                                         resources->remote_mmap_descriptor,
                                         resources->remote_mmap_descriptor_size,
                                         resources->doca_device,
                                         &(resources->remote_mmap));

    if (result != DOCA_SUCCESS) {
        DOCA_LOG_ERR("Failed to create mmap from export: %s", doca_error_get_descr(result));
        return result;
    }
    DOCA_LOG_INFO("\nStarting benchmark: %d transfers\n", NUM_TRANSFERS);
    clock_gettime(CLOCK_MONOTONIC, &resources->start_time);

    resources->transfers_left = NUM_TRANSFERS;

    for (int i = 0; i < 4; i++) {
        if (resources->transfers_left == 0) break;

        result = rdma_write_prepare_and_submit_task(resources);
        if (result != DOCA_SUCCESS) break;

        resources->transfers_left--;
    }
    return result;
}
```

# Buffer offset

- While buffers were writing in the same place in the previous version, now with the pipeline they cannot do that as one would overwrite the other. So, I introduced offset to the memory address each one was writing at.

- The offset I introduced was:  
$$\text{MAX\_BUFF\_SIZE} * (\text{NUM\_TRANSFERS} - \text{resources->transfers\_left})$$

```
/* Add src buffer to DOCA buffer inventory from the remote mmap */
result = doca_buf_inventory_buf_get_by_data(resources->buf_inventory,
                                             resources->mmap,
                                             resources->mmap_memrange + MAX_BUFF_SIZE * (NUM_TRANSFERS - resources->transfers_left),
                                             MAX_BUFF_SIZE,
                                             &resources->src_buf);
if (result != DOCA_SUCCESS) {
    DOCA_LOG_ERR("Failed to allocate DOCA buffer to DOCA buffer inventory: %s",
                 doca_error_get_descr(result));
    return result;
}

/* Set data of src buffer to be the string we want to write */
result = doca_buf_get_data(resources->src_buf, &src_buf_data);
if (result != DOCA_SUCCESS) {
    DOCA_LOG_ERR("Failed to get source buffer data: %s", doca_error_get_descr(result));
    goto destroy_src_buf;
}
// strncpy(src_buf_data, resources->cfg->write_string, write_string_len);
memset(src_buf_data, 'S', MAX_BUFF_SIZE);
/* Add dst buffer to DOCA buffer inventory */
result = doca_buf_inventory_buf_get_by_addr(resources->buf_inventory,
                                             resources->remote_mmap,
                                             remote_mmap_range + MAX_BUFF_SIZE * (NUM_TRANSFERS - resources->transfers_left),
                                             MAX_BUFF_SIZE,
                                             &resources->dst_buf);
if (result != DOCA_SUCCESS) {
    DOCA_LOG_ERR("Failed to allocate DOCA buffer to DOCA buffer inventory: %s",
                 doca_error_get_descr(result));
    goto destroy_src_buf;
}
```

- 
- Also fixed an error in the freeing of the buffers
  - And fixed the premature ending of the benchmark by checking specifically if number of remaining tasks is zero.

# Responder

---

- I added a random offset to the printed 64 bytes so I check if the memory is correctly written.
- Also in BOTH the files I introduced a var signal\_offset that skips the written area in the memory in order for the send/receive buffer to use.

# Times

---

- Times:
  - Throughput: 31.4731 Gbps
  - Throughput: 29.9341 Gbps
  - Throughput: 30.1986 Gbps
  - Throughput: 31.2671 Gbps
  - Throughput: 30.7210 Gbps
- Average = 30.718780
- Standard deviation = 0.662565

# Version 5 (After last meeting)

## Requester

- I filled the memory before hand in order to save time.
  - Drastically increased message size. (All versions now have the same message size for comparison purposes, but I increased the size here to 64MB for the first time.)
  - Also removed DOCA\_LOG\_INFO instructions from the submit write task func because they were wasting time and making excess noise.

```

result = doca_mmap_create_from_export(NULL,
                                      resources->remote_mmap_descriptor,
                                      resources->remote_mmap_descriptor_size,
                                      resources->doca_device,
                                      &(resources->remote_mmap));

if (result != DOCA_SUCCESS) {
    DOCA_LOG_ERR("Failed to create mmap from export: %s", doca_error_get_descr(result));
    return result;
}

for (int j=0;j<NUM_TRANSFERS;j++){
    memset(resources->mmap_memrange + MAX_BUFF_SIZE * j, 'A' + j, MAX_BUFF_SIZE);
}

DOCA_LOG_INFO("\nStarting benchmark: %d transfers\n", NUM_TRANSFERS);
clock_gettime(CLOCK_MONOTONIC, &resources->start_time);

resources->transfers_left = NUM_TRANSFERS;

for (int i = 0; i < PIPELINE_DEPTH; i++) {
    if (resources->transfers_left == 0) break;

    result = rdma_write_prepare_and_submit_task(resources);
    if (result != DOCA_SUCCESS) break;

    resources->transfers_left--;
}

return result;
}

```

# Responder:

---

- I now print the first 64 bytes of each buffer.

```
union doca_data ctx_user_data)

{
    doca_error_t result = DOCA_SUCCESS;
    // char buffer[MAX_BUFF_SIZE+1];
    struct rdma_resources *resources = (struct rdma_resources *)ctx_user_data.ptr;
    DOCA_LOG_INFO("\n Received final signal from requester. Benchmark complete.\n");

    /* The RDMA Write target memory (where requester wrote data) */

    /* Print the first bytes written by the requester */
    char print_buf[65];
    for(int j=0; j<NUM_TRANSFERS;j++){
        char *rdma_written_data = (char *)resources->mmap_memrange + MAX_BUFF_SIZE * j;
        memcpy(print_buf, rdma_written_data, 64);
        print_buf[64] = '\0';
        DOCA_LOG_INFO("Some 64 bytes written via RDMA Write: \"%s\"", print_buf);
    }

    /* Print the received message ("done") separately */
}
```

# Times

---

- Times:
  - Throughput: 64.0851 Gbps
  - Throughput: 48.3647 Gbps
  - Throughput: 46.9563 Gbps
  - Throughput: 46.3695 Gbps
  - Throughput: 43.2281 Gbps
- Average = 49.800740
- Standard deviation = 8.203185

# Version 6:

---

- Requester:
  - Introduced a variable that shows the number of available buffers in memory (NUM\_CHUNKS\_IN\_BUFFER)
  - And made possible to reuse the memory allowing to send more than MEMORY/MAM\_BUFF\_SIZE messages, by circling around.
  - I also fetch remote mmap only once.
- In responder fixed a small issue with freeing the buffer of the receive task.

# Getting remote mmap once

---

```
if (result != DOCA_SUCCESS) {
    DOCA_LOG_ERR("Failed to create mmap from export: %s", doca_error_get_descr(result));
    return result;
}

result = doca_mmap_get_memrange(resources->remote_mmap, (void **)&resources->remote_mmap_range, &resources->remote_mmap_range_len);
if (result != DOCA_SUCCESS) {
    DOCA_LOG_ERR("Failed to get DOCA memory map range: %s", doca_error_get_descr(result));
    return result;
}

for (int j=0;j<NUM_CHUNKS_IN_BUFFER;j++){
```

# Circling logic

---

The image has the code for the local buffer. Its similar for the remote buffer

```
    ...
long long int offset = MAX_BUFF_SIZE * ((NUM_TRANSFERS - resources->transfers_left)%NUM_CHUNKS_IN_BUFFER);

/* Add src buffer to DOCA buffer inventory from the remote mmap */
result = doca_buf_inventory_buf_get_by_data(resources->buf_inventory,
                                             resources->mmap,
                                             resources->mmap_memrange + offset,
                                             MAX_BUFF_SIZE,
                                             &resources->src_buf);

if (result != DOCA_SUCCESS) {
    DOCA_LOG_ERR("Failed to allocate DOCA buffer to DOCA buffer inventory: %s",
                 doca_error_get_descr(result));
    return result;
}
```

# Times

---

- Times:
  - Throughput: 46.5645 Gbps
  - Throughput: 47.4954 Gbps
  - Throughput: 43.8896 Gbps
  - Throughput: 64.3808 Gbps
  - Throughput: 48.4482 Gbps
- Average = 50.155700
- Standard deviation = 8.131835

# Version 7

---

- I increased the size of the buffer inventory and I pre allocated 1000 buffers to not allocate one each time I want to submit a task. This uses a lot more memory and it is fixed that I cannot use memory after the 1000<sup>th</sup> buffer has ended, but it is faster. So, in the key value store application the buffers just have to be the size of the storage.
- To do that I needed to add local\_bufs and remote\_bufs arrays in the resources struct.

No changes to responder

# Buffer pre allocation

```
long long int offset = 0;
for(int i=0; i<NUM_TRANSFERS;i++){
    /* Add src buffer to DOCA buffer inventory from the remote mmap */
    tmp_result = doca_buf_inventory_buf_get_by_data(resources->buf_inventory,
                                                    resources->mmap,
                                                    resources->mmap_memrange + offset,
                                                    MAX_BUFF_SIZE,
                                                    &resources->local_bufs[i]);
    if (tmp_result != DOCA_SUCCESS) {
        DOCA_LOG_ERR("Failed to allocate DOCA buffer to DOCA buffer inventory: %s",
                     doca_error_get_descr(tmp_result));
        // DOCA_ERROR_PROPAGATE(*first_encountered_error, tmp_result);
        return tmp_result;
    }

    tmp_result = doca_buf_inventory_buf_get_by_addr(resources->buf_inventory,
                                                    resources->remote_mmap,
                                                    resources->remote_mmap_range + offset,
                                                    MAX_BUFF_SIZE,
                                                    &resources->remote_bufs[i]);
    if (tmp_result != DOCA_SUCCESS) {
        DOCA_LOG_ERR("Failed to allocate DOCA buffer to DOCA buffer inventory: %s",
                     doca_error_get_descr(tmp_result));
        // DOCA_ERROR_PROPAGATE(*first_encountered_error, tmp_result);
        return tmp_result;
    }
    if (tmp_result == DOCA_SUCCESS) {
        size_t buf_len;
        void *buf_data;
        doca_buf_get_len(resources->remote_bufs[i], &buf_len);
        doca_buf_get_data(resources->remote_bufs[i], &buf_data);
        DOCA_LOG_INFO("Remote buf %d created. Addr: %p, Len: %zu",
                      i, buf_data, buf_len);
    }
    bffset = (offset + MAX_BUFF_SIZE)%PHYSICAL_BUFFER_SIZE;
}
```

```
/* Create DOCA buffer inventory */
result = doca_buf_inventory_create(2*NUM_TRANSFERS+1, &resources.buf_inventory);
if (result != DOCA_SUCCESS) {
    DOCA_LOG_ERR("Failed to create DOCA buffer inventory: %s", doca_error_get_descr(result));
    goto destroy_resources;
}
```

# Freeing buffers in the end

---

```
case DOCA_CTX_STATE_STOPPING:  
    /**  
     * doca_ctx_stop() has been called.  
     * In this sample, this happens either due to a failure encountered, in which case doca_pe_progress()  
     * will cause any inflight task to be flushed, or due to the successful compilation of the sample flow.  
     * In both cases, in this sample, doca_pe_progress() will eventually transition the context to idle  
     * state.  
     */  
    long long int offset = 0;  
    for(int i=0; i<NUM_TRANSFERS;i++){  
        result = doca_buf_dec_refcount(resources->local_bufs[i], NULL);  
        if (result != DOCA_SUCCESS) {  
            DOCA_LOG_ERR("Failed to decrease src_buf count: %s", doca_error_get_descr(result));  
            DOCA_ERROR_PROPAGATE(result, result);  
        }  
  
        result = doca_buf_dec_refcount(resources->remote_bufs[i], NULL);  
        if (result != DOCA_SUCCESS) {  
            DOCA_LOG_ERR("Failed to decrease dst_buf count: %s", doca_error_get_descr(result));  
            DOCA_ERROR_PROPAGATE(result, result);  
        }  
        offset+=MAX_BUFF_SIZE;  
    }  
    DOCA_LOG_INFO("RDMA context entered into stopping state. Any inflight tasks will be flushed sent_tasks = /%d/\n",counter);  
    break;  
case DOCA_CTX_STATE_IDLE:  
    DOCA_LOG_INFO("RDMA context has been stopped");  
  
    /* We can stop progressing the PE */  
    resources->run_pe_progress = false;  
    break;
```

# Times

---

- Times:
  - Throughput: 47.6591 Gbps
  - Throughput: 60.8390 Gbps
  - Throughput: 43.9465 Gbps
  - Throughput: 35.2375 Gbps
  - Throughput: 64.5532 Gbps
  - Throughput: 48.2616 Gbps
- Average = 50.082817
- Standard deviation = 10.884085

# Version 8

---

Pre allocated task pool

```
///////////
for (int i = 0; i < PIPELINE_DEPTH; i++) {
    result = doca_rdma_task_write_allocate_init(resources->rdma,
                                                resources->connections[0],
                                                NULL, // Buffer will be set later
                                                NULL, // Buffer will be set later
                                                (union doca_data){.ptr = &resources->first_encountered_error},
                                                &resources->write_tasks[i]);
    if (result != DOCA_SUCCESS) {
        DOCA_LOG_ERR("Failed to allocate RDMA write task: %s", doca_error_get_descr(result));
        return result;
    }
}
```

# Reuse the pre allocated tasks

```
static doca_error_t rdma_write_prepare_and_submit_task(struct rdma_resources *resources)
{
    // struct doca_rdma_task_write *rdma_write_task = NULL;
    // union doca_data task_user_data = {0};
    doca_error_t result; //, tmp_result;

    int data_idx = (NUM_TRANSFERS - resources->transfers_left);

    struct doca_buf *current_src_buf = resources->local_bufs[data_idx];
    struct doca_buf *current_dst_buf = resources->remote_bufs[data_idx];

    int task_idx = data_idx % PIPELINE_DEPTH;
    struct doca_rdma_task_write *task = resources->write_tasks[task_idx];
    doca_rdma_task_write_set_src_buf(task, current_src_buf);
    doca_rdma_task_write_set_dst_buf(task, current_dst_buf);

    /* Submit RDMA write task */
    // DOCA_LOG_INFO("Submitting RDMA write task that writes SSS... to the responder");
    resources->num_remaining_tasks++;
    result = doca_task_submit(doca_rdma_task_write_as_task(task));
    if (result != DOCA_SUCCESS) {
        DOCA_LOG_ERR("Failed to submit RDMA write task: %s", doca_error_get_descr(result));
        resources->num_remaining_tasks--;
    }

    return result;

    // free_task:
    // doca_task_free(doca_rdma_task_write_as_task(rdma_write_task));
    // return result;
}
```

# Task cleanup

---

```
case DOCA_CTX_STATE_STOPPING:  
    /**  
     * doca_ctx_stop() has been called.  
     * In this sample, this happens either due to a failure encountered, in which case doca_pe_progress()  
     * will cause any inflight task to be flushed, or due to the successful compilation of the sample flow.  
     * In both cases, in this sample, doca_pe_progress() will eventually transition the context to idle  
     * state.  
     */  
    long long int offset = 0;  
    for(int i=0; i<NUM_TRANSFERS;i++){  
        tmp_result = doca_buf_dec_refcount(resources->local_bufs[i], NULL);  
        if (tmp_result != DOCA_SUCCESS) {  
            DOCA_LOG_ERR("Failed to decrease src_buf count: %s", doca_error_get_descr(tmp_result));  
            DOCA_ERROR_PROPAGATE(result, tmp_result);  
        }  
  
        tmp_result = doca_buf_dec_refcount(resources->remote_bufs[i], NULL);  
        if (tmp_result != DOCA_SUCCESS) {  
            DOCA_LOG_ERR("Failed to decrease dst_buf count: %s", doca_error_get_descr(tmp_result));  
            DOCA_ERROR_PROPAGATE(result, tmp_result);  
        }  
        offset+=MAX_BUFF_SIZE;  
    }  
    for (int i = 0; i < PIPELINE_DEPTH; i++) {  
        if (resources->write_tasks[i] != NULL) {  
            doca_task_free(doca_rdma_task_write_as_task(resources->write_tasks[i]));  
        }  
    }  
    DOCA_LOG_INFO("RDMA context entered into stopping state. Any inflight tasks will be flushed sent_tasks = /%d/\n",counter);  
    break;  
case DOCA_CTX_STATE_IDLE:
```

# Times

---

- Times:
  - Throughput: 90.5575 Gbps
  - Throughput: 61.5395 Gbps
  - Throughput: 67.5940 Gbps
  - Throughput: 91.4312 Gbps
  - Throughput: 64.4800 Gbps
  - Throughput: 94.9797 Gbps
- Average = 78.430317
- Standard deviation = 15.409805

# lb\_write\_bw max throughput test

---

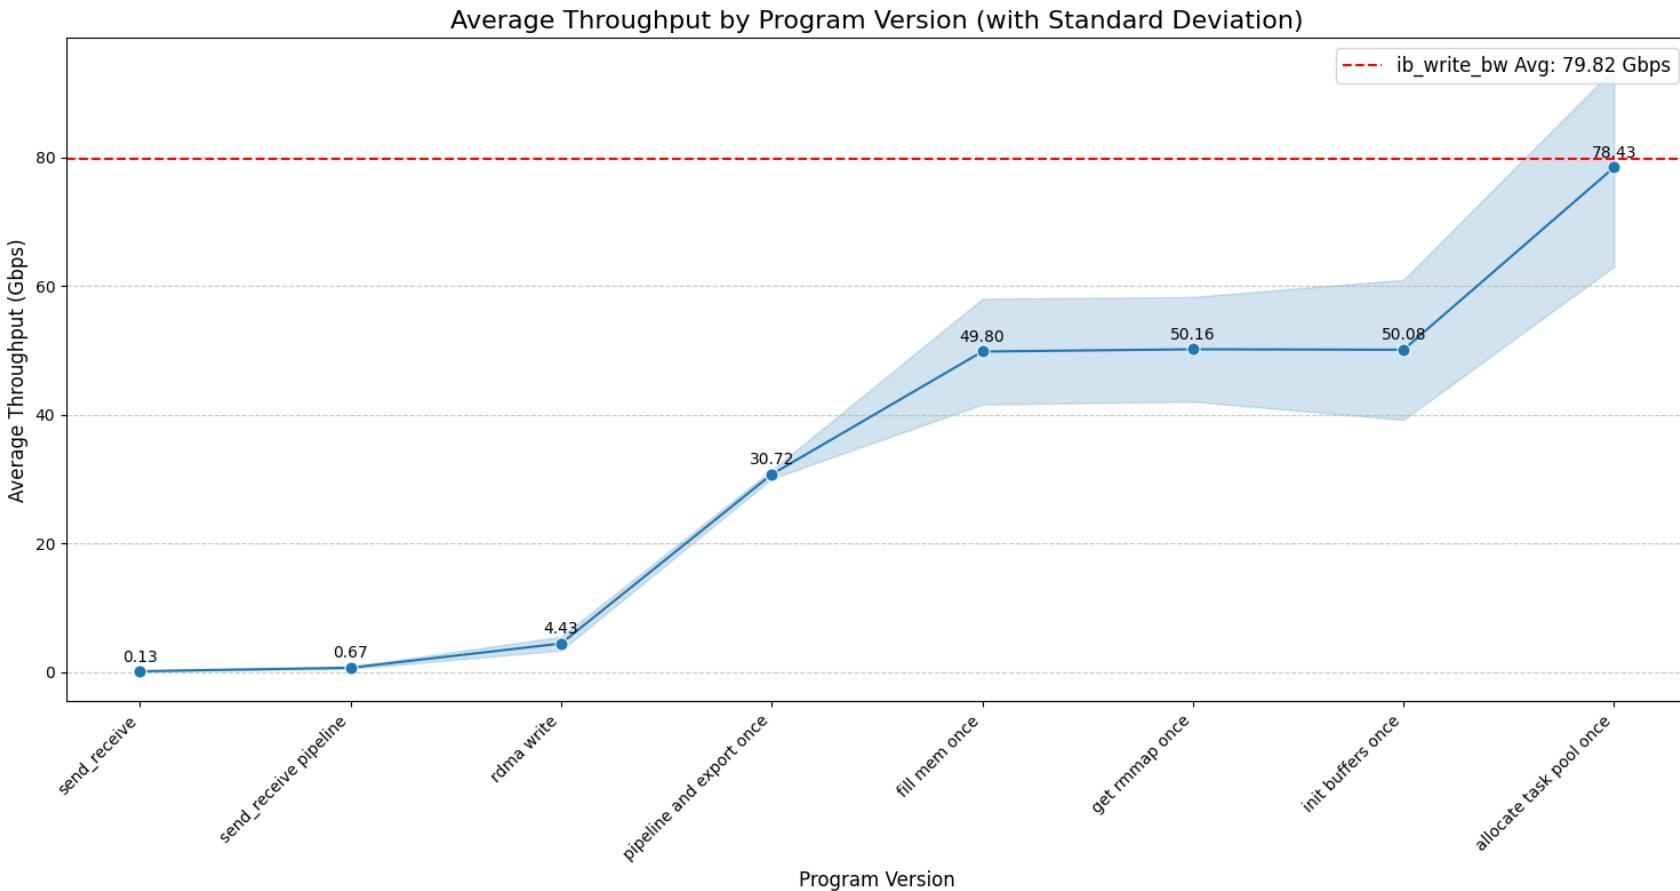
- To run the test run:
  - ib\_write\_bw, instruction to the server side
  - ib\_write\_bw <server\_ip>, instruction to the client side

# lb\_write\_bw max throughput test times

---

- The test indicated that the CPU was in a power-saving state so it could not reach its maximum potential.
- Times:
  - Throughput: 88.24 Gbps
  - Throughput: 89.87 Gbps
  - Throughput: 72.26 Gbps
  - Throughput: 69.93 Gbps
  - Throughput: 78.63 Gbps
  - Throughput: 80.0 Gbps
- Average = 79.821667
- Standard deviation = 8.103247

# Times comparison



---

During the development of these versions some changes had to be made in both:

- `rdma_common.c`
- `rdma_common.h`

# Changes in rdma\_common.h

---

```
#define MEM_RANGE_LEN (1LL * 1024 * 1024 * 1024)

struct rdma_resources {
    // ... existing members

    struct timespec start_time;
    struct timespec end_time;
    int transfers_left;
    char *remote_mmap_range;
    size_t remote_mmap_range_len;
    struct doca_buf *local_bufs[1000];
    struct doca_buf *remote_bufs[1000];
    int cur_buf_idx;
    struct doca_rdma_task_write *write_tasks[128];
};
```

# Changes in rdma\_common.c

---

- Problem:

While developing the benchmarks, the DOCA driver gave a performance warning that memory range isn't aligned to 64B.

The problem was in the allocate\_rdma\_resources function, which was using calloc to get memory for the RDMA buffers.

The standard calloc (and malloc) functions do not guarantee 64-byte alignment, so they sometimes returned an unaligned address, triggering the warning.

# Solution:

---

- I used posix\_memalign instead of calloc.

```
void *memrange_ptr = NULL;  
int ret = posix_memalign(&memrange_ptr, 64, MEM_RANGE_LEN);
```